

deep learning CNN Data augmentation to address overfitting

May 11, 2023

```
[5]: import matplotlib.pyplot as plt
import numpy as np
import cv2
import os
import PIL
import tensorflow as tf

from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.models import Sequential
```

```
[6]: !pip install cv2
```

```
ERROR: Could not find a version that satisfies the requirement cv2 (from
versions: none)
ERROR: No matching distribution found for cv2
```

```
[3]: !pip install opencv-python
```

```
Collecting opencv-python
  Downloading opencv_python-4.7.0.72-cp37-abi3-win_amd64.whl (38.2 MB)
Requirement already satisfied: numpy>=1.19.3 in
c:\users\rakesh\anaconda3\lib\site-packages (from opencv-python) (1.21.5)
Installing collected packages: opencv-python
Successfully installed opencv-python-4.7.0.72
```

```
[4]: import cv2
```

```
[7]: import PIL
```

```
[12]: dataset_url="https://storage.googleapis.com/download.tensorflow.org/
↳example_images/flower_photos.tgz"
data_dir=tf.keras.utils.get_file('flower_photos',origin=dataset_url,
↳cache_dir='.',untar=True)
```

```
Downloading data from https://storage.googleapis.com/download.tensorflow.org/exa
mple_images/flower_photos.tgz
228813984/228813984 [=====] - 24s 0us/step
```

```
[13]: data_dir
```

```
[13]: '.\\datasets\\flower_photos'
```

```
[14]: import pathlib
data_dir=pathlib.Path(data_dir)
data_dir
```

```
[14]: WindowsPath('datasets/flower_photos')
```

```
[15]: list(data_dir.glob('*/*.jpg'))[:5]
```

```
[15]: [WindowsPath('datasets/flower_photos/daisy/100080576_f52e8ee070_n.jpg'),
WindowsPath('datasets/flower_photos/daisy/10140303196_b88d3d6cec.jpg'),
WindowsPath('datasets/flower_photos/daisy/10172379554_b296050f82_n.jpg'),
WindowsPath('datasets/flower_photos/daisy/10172567486_2748826a8b.jpg'),
WindowsPath('datasets/flower_photos/daisy/10172636503_21bededa75_n.jpg')]
```

```
[17]: image_count=len(list(data_dir.glob('*/*.jpg')))
image_count
```

```
[17]: 3670
```

```
[18]: PIL.Image.open(str(roses[1]))
```

```
-----
NameError                                Traceback (most recent call last)
Input In [18], in <cell line: 1>()
----> 1 PIL.Image.open(str(roses[1]))

NameError: name 'roses' is not defined
```

```
[19]: roses=list(data_dir.glob('roses/*'))
roses[:5]
```

```
[19]: [WindowsPath('datasets/flower_photos/roses/10090824183_d02c613f10_m.jpg'),
WindowsPath('datasets/flower_photos/roses/102501987_3cdb8e5394_n.jpg'),
WindowsPath('datasets/flower_photos/roses/10503217854_e66a804309.jpg'),
WindowsPath('datasets/flower_photos/roses/10894627425_ec76bbc757_n.jpg'),
WindowsPath('datasets/flower_photos/roses/110472418_87b6a3aa98_m.jpg')]
```

```
[85]: PIL.Image.open(str(roses[2]))
```

```
[85]:
```



```
[25]: tulips=list(data_dir.glob('tulips/*'))  
tulips[:5]
```

```
[25]: [WindowsPath('datasets/flower_photos/tulips/100930342_92e8746431_n.jpg'),  
WindowsPath('datasets/flower_photos/tulips/10094729603_eeca3f2cb6.jpg'),  
WindowsPath('datasets/flower_photos/tulips/10094731133_94a942463c.jpg'),  
WindowsPath('datasets/flower_photos/tulips/10128546863_8de70c610d.jpg'),  
WindowsPath('datasets/flower_photos/tulips/10163955604_ae0b830975_n.jpg')]
```

```
[29]: PIL.Image.open(str(tulips[0]))
```

```
[29]:
```



Read flowers images from disk into numpy array using opencv

```
[31]: flowers_images_dict={
    'roses':list(data_dir.glob('roses/*')),
    'daisy': list(data_dir.glob('daisy/*')),
    'dandelion': list(data_dir.glob('dandelion/*')),
    'sunflowers': list(data_dir.glob('sunflowers/*')),
    'tulips': list(data_dir.glob('tulips/*')),
}
```

```
[32]: flower_label_dict={
    'roses':0,
    'daisy': 1,
    'dandelion': 2,
    'sunflowers': 3,
    'tulips': 4,
}
```

```
[33]: flowers_images_dict['roses'][:5]
```

```
[33]: [WindowsPath('datasets/flower_photos/roses/10090824183_d02c613f10_m.jpg'),
WindowsPath('datasets/flower_photos/roses/102501987_3cdb8e5394_n.jpg'),
WindowsPath('datasets/flower_photos/roses/10503217854_e66a804309.jpg'),
WindowsPath('datasets/flower_photos/roses/10894627425_ec76bbc757_n.jpg'),
WindowsPath('datasets/flower_photos/roses/110472418_87b6a3aa98_m.jpg')]
```

```
[34]: str(flowers_images_dict['roses'][0])
```

```
[34]: 'datasets\\flower_photos\\roses\\10090824183_d02c613f10_m.jpg'
```

```
[35]: img=cv2.imread(str(flowers_images_dict['roses'][0]))
```

```
[36]: img.shape
```

```
[36]: (240, 179, 3)
```

```
[37]: cv2.resize(img,(180,180)).shape
```

```
[37]: (180, 180, 3)
```

```
[40]: x,y=[],[]  
for flower_name, images in flowers_images_dict.items():  
    for image in images:  
        img=cv2.imread(str(image))  
        resized_images=cv2.resize(img,(180,180))  
        x.append(resized_images)  
        y.append(flower_label_dict[flower_name])
```

```
[41]: x=np.array(x)  
y=np.array(y)
```

Train test split

```
[42]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,random_state=0)
```

Preprocessing: scale images

```
[44]: x_train_scaled=x_train/255  
x_test_scaled=x_test/255
```

Build convolutional neural network and train it

```
[51]: num_classes=5  
  
model=Sequential([  
    layers.Conv2D(16,3,padding='same',activation='relu'),  
    layers.MaxPooling2D(),  
    layers.Conv2D(16,3,padding='same',activation='relu'),  
    layers.MaxPooling2D(),  
    layers.Conv2D(16,3,padding='same',activation='relu'),  
    layers.MaxPooling2D(),  
  
    layers.Flatten(),  
    layers.Dense(128,activation='relu'),  
    layers.Dense(num_classes)  
)
```

```

model.compile(optimizer='adam',
              loss=tf.keras.losses.
↳SparseCategoricalCrossentropy(from_logits=True),
              metrics=['accuracy']
              )

model.fit(x_train_scaled,y_train,epochs=30)

```

```

Epoch 1/30
86/86 [=====] - 37s 415ms/step - loss: 1.3714 -
accuracy: 0.4033
Epoch 2/30
86/86 [=====] - 36s 418ms/step - loss: 1.0256 -
accuracy: 0.6017
Epoch 3/30
86/86 [=====] - 34s 396ms/step - loss: 0.8632 -
accuracy: 0.6686
Epoch 4/30
86/86 [=====] - 34s 392ms/step - loss: 0.7067 -
accuracy: 0.7344
Epoch 5/30
86/86 [=====] - 34s 397ms/step - loss: 0.5470 -
accuracy: 0.8063
Epoch 6/30
86/86 [=====] - 34s 393ms/step - loss: 0.4038 -
accuracy: 0.8616
Epoch 7/30
86/86 [=====] - 34s 392ms/step - loss: 0.2449 -
accuracy: 0.9182
Epoch 8/30
86/86 [=====] - 34s 392ms/step - loss: 0.1765 -
accuracy: 0.9455
Epoch 9/30
86/86 [=====] - 34s 394ms/step - loss: 0.1035 -
accuracy: 0.9677
Epoch 10/30
86/86 [=====] - 34s 396ms/step - loss: 0.0758 -
accuracy: 0.9811
Epoch 11/30
86/86 [=====] - 36s 413ms/step - loss: 0.0702 -
accuracy: 0.9833
Epoch 12/30
86/86 [=====] - 35s 402ms/step - loss: 0.0434 -
accuracy: 0.9873
Epoch 13/30
86/86 [=====] - 34s 391ms/step - loss: 0.0378 -
accuracy: 0.9916

```

Epoch 14/30
86/86 [=====] - 34s 397ms/step - loss: 0.0170 -
accuracy: 0.9967
Epoch 15/30
86/86 [=====] - 34s 392ms/step - loss: 0.0246 -
accuracy: 0.9949
Epoch 16/30
86/86 [=====] - 34s 398ms/step - loss: 0.0276 -
accuracy: 0.9949
Epoch 17/30
86/86 [=====] - 34s 390ms/step - loss: 0.0134 -
accuracy: 0.9978
Epoch 18/30
86/86 [=====] - 34s 397ms/step - loss: 0.1131 -
accuracy: 0.9651
Epoch 19/30
86/86 [=====] - 34s 391ms/step - loss: 0.0653 -
accuracy: 0.9797
Epoch 20/30
86/86 [=====] - 34s 396ms/step - loss: 0.0319 -
accuracy: 0.9927
Epoch 21/30
86/86 [=====] - 33s 381ms/step - loss: 0.0097 -
accuracy: 0.9989
Epoch 22/30
86/86 [=====] - 34s 391ms/step - loss: 0.0124 -
accuracy: 0.9975
Epoch 23/30
86/86 [=====] - 33s 382ms/step - loss: 0.0113 -
accuracy: 0.9982
Epoch 24/30
86/86 [=====] - 34s 390ms/step - loss: 0.0082 -
accuracy: 0.9993
Epoch 25/30
86/86 [=====] - 33s 382ms/step - loss: 0.0041 -
accuracy: 0.9993
Epoch 26/30
86/86 [=====] - 36s 415ms/step - loss: 0.0057 -
accuracy: 0.9993
Epoch 27/30
86/86 [=====] - 33s 387ms/step - loss: 0.0019 -
accuracy: 0.9996
Epoch 28/30
86/86 [=====] - 34s 390ms/step - loss: 0.0032 -
accuracy: 0.9996
Epoch 29/30
86/86 [=====] - 33s 381ms/step - loss: 0.0025 -
accuracy: 0.9996

```
Epoch 30/30
86/86 [=====] - 34s 397ms/step - loss: 0.0058 -
accuracy: 0.9993
```

```
[51]: <keras.callbacks.History at 0x1f2aba71a00>
```

```
[53]: model.evaluate(x_test_scaled,y_test)
```

```
29/29 [=====] - 4s 125ms/step - loss: 2.8721 -
accuracy: 0.6264
```

```
[53]: [2.8721227645874023, 0.6263616681098938]
```

Here we see that while train accuracy is very high (99%), the test accuracy is significantly low (66.99%) indicating overfitting. Let's make some predictions before we use data augmentation to address overfitting

```
[55]: predictions=model.predict(x_test_scaled)
      predictions
```

```
29/29 [=====] - 4s 137ms/step
```

```
[55]: array([[ 8.336534 , 15.13717 ,  5.4003415, -10.471546 ,
            -8.57661  ],
            [ 8.395139 , -2.4289827, -6.2180243 ,  1.9653721 ,
             5.3486357 ],
            [-8.2515745 ,  2.6009653 , 15.647339 , -11.676489 ,
            -5.8864093 ],
            ...,
            [-4.237256 ,  5.8607697 , -1.9202993 ,  6.85216 ,
            -0.16207433],
            [-3.3413472 , -0.15840845,  3.7371826 , -3.348835 ,
             7.909586  ],
            [ 3.9100323 , -13.481786 ,  2.1709642 ,  9.328533 ,
             5.725955  ]], dtype=float32)
```

```
[57]: score=tf.nn.softmax(predictions[0])
```

```
[58]: np.argmax(score)
```

```
[58]: 1
```

```
[59]: y_test[0]
```

```
[59]: 1
```


0.0.1 Improve Test Accuracy Using Data Augmentation

```
[75]: data_augmentation=keras.Sequential([
        layers.experimental.preprocessing.RandomFlip("horizontal",
                                                    input_shape=(img_height,
                                                                    img_width,
                                                                    3)),

        layers.experimental.preprocessing.RandomRotation(0.1),
        layers.experimental.preprocessing.RandomZoom(0.9)
    ])
```

```
[71]: plt.axis('off')
      plt.imshow(x[0])
```

```
[71]: <matplotlib.image.AxesImage at 0x1f1067fcc70>
```



```
[77]: plt.axis('off')
      plt.imshow(data_augmentation(x)[0].numpy().astype("uint8"))
```

WARNING:tensorflow:Using a while_loop for converting RngReadAndSkip cause there is no registered converter for this op.

WARNING:tensorflow:Using a while_loop for converting Bitcast cause there is no registered converter for this op.

WARNING:tensorflow:Using a while_loop for converting Bitcast cause there is no registered converter for this op.

WARNING:tensorflow:Using a while_loop for converting StatelessRandomUniformV2 cause there is no registered converter for this op.

WARNING:tensorflow:Using a while_loop for converting ImageProjectiveTransformV3 cause there is no registered converter for this op.
WARNING:tensorflow:Using a while_loop for converting RngReadAndSkip cause there is no registered converter for this op.
WARNING:tensorflow:Using a while_loop for converting Bitcast cause there is no registered converter for this op.
WARNING:tensorflow:Using a while_loop for converting Bitcast cause there is no registered converter for this op.
WARNING:tensorflow:Using a while_loop for converting StatelessRandomUniformV2 cause there is no registered converter for this op.
WARNING:tensorflow:Using a while_loop for converting ImageProjectiveTransformV3 cause there is no registered converter for this op.

[77]: <matplotlib.image.AxesImage at 0x1f106d46c10>



```
[82]: num_classes=5

model1=Sequential([
    data_augmentation,
    layers.Conv2D(18,3,padding='same',activation='relu'),
    layers.MaxPooling2D(),

    layers.Conv2D(32,3,padding='same',activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(64,3,padding='same',activation='relu'),
    layers.MaxPooling2D(),
```

```

        layers.Dropout(0.2),
        layers.Flatten(),

        layers.Dense(128,activation='relu')    ,
        layers.Dense(num_classes)
    ])

model1.compile(optimizer='adam',
               loss=tf.keras.losses.
↳SparseCategoricalCrossentropy(from_logits=True),
               metrics=['accuracy'])

model.fit(x_train_scaled,y_train,epochs=30)

```

```

Epoch 1/30
86/86 [=====] - 40s 461ms/step - loss: 9.3052e-04 -
accuracy: 0.9996
Epoch 2/30
86/86 [=====] - 49s 569ms/step - loss: 0.0013 -
accuracy: 0.9993
Epoch 3/30
86/86 [=====] - 42s 487ms/step - loss: 9.2269e-04 -
accuracy: 0.9996
Epoch 4/30
86/86 [=====] - 34s 398ms/step - loss: 0.0030 -
accuracy: 0.9993
Epoch 5/30
86/86 [=====] - 34s 398ms/step - loss: 0.4573 -
accuracy: 0.8583
Epoch 6/30
86/86 [=====] - 35s 412ms/step - loss: 0.1017 -
accuracy: 0.9709
Epoch 7/30
86/86 [=====] - 33s 388ms/step - loss: 0.0358 -
accuracy: 0.9887
Epoch 8/30
86/86 [=====] - 34s 393ms/step - loss: 0.0102 -
accuracy: 0.9978
Epoch 9/30
86/86 [=====] - 33s 387ms/step - loss: 0.0033 -
accuracy: 0.9996
Epoch 10/30
86/86 [=====] - 35s 403ms/step - loss: 0.0025 -
accuracy: 0.9996
Epoch 11/30
86/86 [=====] - 34s 401ms/step - loss: 0.0018 -

```

```

accuracy: 0.9996
Epoch 12/30
86/86 [=====] - 34s 396ms/step - loss: 0.0021 -
accuracy: 0.9993
Epoch 13/30
86/86 [=====] - 33s 386ms/step - loss: 0.0019 -
accuracy: 0.9993
Epoch 14/30
86/86 [=====] - 34s 392ms/step - loss: 0.0015 -
accuracy: 0.9996
Epoch 15/30
86/86 [=====] - 34s 390ms/step - loss: 0.0014 -
accuracy: 0.9996
Epoch 16/30
86/86 [=====] - 34s 391ms/step - loss: 0.0016 -
accuracy: 0.9996
Epoch 17/30
86/86 [=====] - 33s 387ms/step - loss: 0.0014 -
accuracy: 0.9996
Epoch 18/30
86/86 [=====] - 34s 391ms/step - loss: 0.0012 -
accuracy: 0.9996
Epoch 19/30
86/86 [=====] - 34s 390ms/step - loss: 0.0012 -
accuracy: 0.9996
Epoch 20/30
86/86 [=====] - 34s 401ms/step - loss: 0.0012 -
accuracy: 0.9996
Epoch 21/30
86/86 [=====] - 34s 394ms/step - loss: 0.0021 -
accuracy: 0.9993
Epoch 22/30
86/86 [=====] - 34s 394ms/step - loss: 0.0012 -
accuracy: 0.9996
Epoch 23/30
86/86 [=====] - 33s 388ms/step - loss: 0.0012 -
accuracy: 0.9996
Epoch 24/30
86/86 [=====] - 34s 393ms/step - loss: 0.0015 -
accuracy: 0.9993
Epoch 25/30
86/86 [=====] - 33s 387ms/step - loss: 0.0011 -
accuracy: 0.9996
Epoch 26/30
86/86 [=====] - 34s 394ms/step - loss: 8.7374e-04 -
accuracy: 0.9996
Epoch 27/30
86/86 [=====] - 33s 388ms/step - loss: 0.0011 -

```

```
accuracy: 0.9996
Epoch 28/30
86/86 [=====] - 34s 395ms/step - loss: 0.0015 -
accuracy: 0.9993
Epoch 29/30
86/86 [=====] - 33s 388ms/step - loss: 0.0013 -
accuracy: 0.9993
Epoch 30/30
86/86 [=====] - 34s 394ms/step - loss: 0.0012 -
accuracy: 0.9993
```

```
[82]: <keras.callbacks.History at 0x1f25e6182b0>
```

```
[83]: model1.evaluate(x_test,y_test)
```

```
29/29 [=====] - 6s 167ms/step - loss: 26.0198 -
accuracy: 0.1438
```

```
[83]: [26.019821166992188, 0.1437908560037613]
```

```
[ ]:
```