

# DEEP LEARNING CNN FOR IMG

May 11, 2023

```
[1]: import tensorflow as tf
      from tensorflow.keras import datasets, layers, models
      import matplotlib.pyplot as plt
      import numpy as np
```

## 1 Load the dataset

```
[14]: (x_train,y_train),(x_test,y_test)= datasets.cifar10.load_data()
      x_train.shape
```

```
[14]: (50000, 32, 32, 3)
```

```
[15]: x_test.shape
```

```
[15]: (10000, 32, 32, 3)
```

```
[16]: y_train.shape
```

```
[16]: (50000, 1)
```

```
[17]: y_train[:5]
```

```
[17]: array([[6],
        [9],
        [9],
        [4],
        [1]], dtype=uint8)
```

```
[18]: y_train=y_train.reshape(-1)
      y_train[:5]
```

```
[18]: array([6, 9, 9, 4, 1], dtype=uint8)
```


```
[20]: y_test=y_test.reshape(-1)
      y_test[:5]
```

```
[20]: array([3, 8, 8, 0, 6], dtype=uint8)
```

```
[21]: classes =  
      ↪ ["airplane", "automobile", "bird", "cat", "deer", "dog", "frog", "horse", "ship", "truck"]
```


```
[26]: plt.imshow(x_train[1])
```

```
[26]: <matplotlib.image.AxesImage at 0x21ff837d580>
```



```
[36]: def plot_sample(x, y, index):  
      plt.figure(figsize = (15,2))  
      plt.imshow(x[index])  
      plt.xlabel(classes[y[index]])
```

```
[38]: plot_sample(x_train, y_train,0)
```



### Normalizing the training data

```
[41]: x_train = x_train / 255.0
      x_test = x_test / 255.0
```

### Build simple artificial neural network for image classification

```
[46]: ann=models.Sequential([
      layers.Flatten(input_shape=(32, 32, 3)),
      layers.Dense(3000,activation='relu'),
      layers.Dense(1000,activation='relu'),
      layers.Dense(10,activation='softmax')

    ])
    ann.compile(
        optimizer='SGD',
        loss='sparse_categorical_crossentropy',
        metrics=['accuracy']
    )

    ann.fit(x_train,y_train,epochs=5)
```

```
Epoch 1/5
1563/1563 [=====] - 108s 68ms/step - loss: 1.8119 -
accuracy: 0.3535
Epoch 2/5
1563/1563 [=====] - 116s 74ms/step - loss: 1.6212 -
accuracy: 0.4289
Epoch 3/5
1563/1563 [=====] - 112s 72ms/step - loss: 1.5388 -
accuracy: 0.4604
Epoch 4/5
1563/1563 [=====] - 114s 73ms/step - loss: 1.4784 -
accuracy: 0.4816
Epoch 5/5
1563/1563 [=====] - 113s 72ms/step - loss: 1.4305 -
accuracy: 0.4958
```

```
[46]: <keras.callbacks.History at 0x21ff9e58d00>
```

```
[ ]: from sklearn.metrics import confusion_matrix, classification_report
     import numpy as np
     y_pred=ann.pred(x_train)
     y_pred_classes=[np.argmax(i) for i in y_pred]
     print('classification report:\n',classification_report(y_test,y_pred_classes))
```

Now let us build a convolutional neural network to train our images

```
[50]: cnn=models.Sequential([
    layers.
    ↪Conv2D(filters=32,kernel_size=(3,3),activation='relu',input_shape=(32, 32,1
    ↪3)),
    layers.MaxPooling2D(2,2),
    layers.
    ↪Conv2D(filters=32,kernel_size=(3,3),activation='relu',input_shape=(32, 32,1
    ↪3)),
    layers.MaxPooling2D(2,2),

    layers.Flatten(),
    layers.Dense(64,activation='relu'),
    layers.Dense(10,activation='softmax')

])
```

```
[51]: cnn.compile(optimizer='adam',
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])
```

```
[52]: cnn.fit(x_train, y_train, epochs=10)
```

```
Epoch 1/10
1563/1563 [=====] - 55s 34ms/step - loss: 1.4824 -
accuracy: 0.4654
Epoch 2/10
1563/1563 [=====] - 46s 30ms/step - loss: 1.1434 -
accuracy: 0.5986
Epoch 3/10
1563/1563 [=====] - 46s 30ms/step - loss: 1.0173 -
accuracy: 0.6467
Epoch 4/10
1563/1563 [=====] - 47s 30ms/step - loss: 0.9416 -
accuracy: 0.6740
Epoch 5/10
1563/1563 [=====] - 46s 30ms/step - loss: 0.8898 -
accuracy: 0.6895
Epoch 6/10
1563/1563 [=====] - 47s 30ms/step - loss: 0.8412 -
accuracy: 0.7092
Epoch 7/10
1563/1563 [=====] - 47s 30ms/step - loss: 0.8042 -
accuracy: 0.7204
Epoch 8/10
1563/1563 [=====] - 47s 30ms/step - loss: 0.7729 -
accuracy: 0.7310
Epoch 9/10
1563/1563 [=====] - 47s 30ms/step - loss: 0.7413 -
```

```
accuracy: 0.7417
Epoch 10/10
1563/1563 [=====] - 46s 29ms/step - loss: 0.7185 -
accuracy: 0.7508
```

```
[52]: <keras.callbacks.History at 0x21ff44dc970>
```

```
[53]: cnn.evaluate(x_test,y_test)
```

```
313/313 [=====] - 4s 10ms/step - loss: 0.9174 -
accuracy: 0.6870
```

```
[53]: [0.9173777103424072, 0.6869999766349792]
```

```
[54]: y_pred = cnn.predict(x_test)
      y_pred[:5]
```

```
313/313 [=====] - 3s 10ms/step
```

```
[54]: array([[1.9725633e-03, 1.5894047e-03, 3.4172602e-02, 6.2655628e-01,
            1.1393116e-03, 2.4108025e-01, 5.3924888e-02, 7.4508460e-04,
            3.8541365e-02, 2.7816434e-04],
            [7.5108828e-03, 6.8859190e-01, 8.8441222e-05, 3.8812198e-07,
            1.3799688e-06, 3.3285879e-08, 2.4247137e-07, 2.9485159e-06,
            2.9701161e-01, 6.7921886e-03],
            [2.8895989e-02, 2.9311344e-01, 3.5146768e-03, 3.6700168e-03,
            7.7218062e-04, 8.3693897e-04, 3.4413037e-03, 1.4111316e-03,
            6.3737273e-01, 2.6971566e-02],
            [9.6753818e-01, 5.6178086e-03, 1.9537913e-02, 7.7568390e-04,
            5.4022152e-04, 4.8479553e-05, 1.6037687e-03, 1.6161461e-03,
            2.6224009e-03, 9.9434881e-05],
            [3.0197575e-06, 8.8355891e-06, 3.1184141e-02, 3.6131172e-03,
            6.2262379e-02, 4.2312237e-04, 9.0248674e-01, 3.6710088e-07,
            1.8193814e-05, 2.3719113e-07]], dtype=float32)
```


```
[55]: y_classes = [np.argmax(element) for element in y_pred]
      y_classes[:5]
```

```
[55]: [3, 1, 8, 0, 6]
```

```
[56]: y_test[:5]
```

```
[56]: array([3, 8, 8, 0, 6], dtype=uint8)
```

```
[57]: plot_sample(x_test, y_test,3)
```



```
[58]: classes[y_classes[3]]
```

```
[58]: 'airplane'
```

```
[ ]:
```